

LINEAR REGRESSION

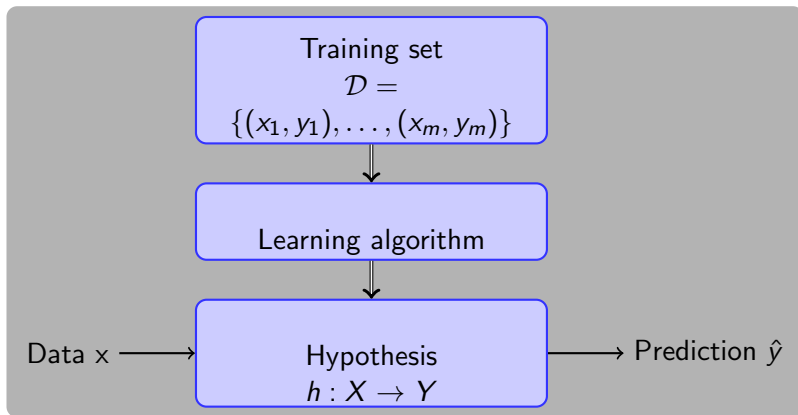
Machine Learning for Autonomous Robots

Dr. Alexander Fabisch
DFKI, Robotics Innovation Center

22 November, 2022 – Bremen, Deutschland

Problem statement


Recap: Supervised Learning



- ▶ If Y is a discrete domain: classification
- ▶ If Y is a continuous domain: regression




An example: Robot trajectory

- ▶ Consider the task of estimating (a posteriori) the actual state of a robot at time t when only noisy measurements exist
- ▶ When no model of the system's dynamics is known: One could try to infer the true trajectory from the data
- ▶ Example data:



Simple approaches

- ▶ Interpolation: Cannot reduce influence of noise
- ▶ Nearest neighbor: Cannot reduce influence of noise
- ▶ Moving window average
 - ▶ How large should moving window be chosen?
 - ▶ Can this size be chosen globally?
 - ▶ How can this be extended to $X = \mathbb{R}^n$?



Linear Regression

Model: Linear function plus gaussian noise

- ▶ Assumption: True, unknown function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is linear in x : $f(x) = \mathbf{w}^T x + \mathbf{b}$, where \mathbf{w} is a n -dimensional vector and \mathbf{b} is a scalar (bias).
- ▶ This can be more conveniently written as $f(x) = \tilde{\mathbf{w}}^T \tilde{x}$ with $\tilde{\mathbf{w}} = (b, w_0, \dots, w_n)^T$ and $\tilde{x} = (1, x_0, \dots, x_n)^T$
- ▶ The measured data is then given by $y^{(i)} = \mathbf{w}^T x^{(i)} + \epsilon^{(i)}$ where $\epsilon^{(i)}$ is the noise in the i -th measurement.
- ▶ It is often assumed that the $\epsilon^{(i)}$ are normally distributed with mean 0 and unknown variance σ^2 : $\epsilon^{(i)} \sim \mathcal{N}(0, \sigma^2)$

Maximum likelihood estimate (1/2)

- ▶ We would like to find a parameter vector $\hat{\mathbf{w}}$ that is as close as possible to the true, unknown \mathbf{w} .
- ▶ We have two sources of information regarding \mathbf{w} :
 - ▶ Our prior belief about \mathbf{w} : $p(\mathbf{w})$.
 - ▶ The observed data $\mathcal{D} = \{(x^{(0)}, y^{(0)}) , \dots , (x^{(m)}, y^{(m)})\}$ and the likelihood of a given w for this data: $p(\mathcal{D}|\mathbf{w})$.
- ▶ For uniform priors we can choose $\hat{\mathbf{w}}$ so that it maximizes the likelihood:
$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w} \in \mathbb{R}^n} p(\mathcal{D}|\mathbf{w})$$
- ▶ When assuming that the elements of \mathcal{D} are independent:
$$\hat{\mathbf{w}} = \arg \max_{\mathbf{w} \in \mathbb{R}^n} \prod_{i=0}^m p((x^{(i)}, y^{(i)})|\mathbf{w})$$


Maximum likelihood estimate (2/2)

- ▶ We assumed $y^{(i)} = \mathbf{w}^T x^{(i)} + \epsilon^{(i)}$ with $\epsilon^{(i)} \sim \mathcal{N}(0, \sigma^2)$.
- ▶ It can be shown that finding w that maximizes $\prod_{i=0}^m p(y^{(i)} | x^{(i)}, w)$ (MLE), is equivalent to finding w that minimizes $\sum_{i=0}^m (y^{(i)} - w^T x^{(i)})^2$ (SSE).
- ▶ This is called sum of squared errors (SSE)
- ▶ Linear **least squares**: $\hat{w} = \arg \min_{w \in \mathbb{R}^n} SSE(w)$
- ▶ How can one determine \hat{w} ? \Rightarrow Approach: *gradient descent*

Gradient descent

- ▶ Gradient descent is an **iterative** method that is based on the intuition that one finally reaches the bottom of a valley when always going downwards on a hill.
- ▶ Mathematical background: Negative gradient is direction of largest decrease of function value!
- ▶ Given objective function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ that is differentiable. Objective: Find minimum (or maximum).
- ▶ Start with some arbitrary $x_0 \in \mathbb{R}^n$.
- ▶ Iteratively update: $x \leftarrow x - \alpha \nabla_x f(x)$
 - ▶ $\nabla_x f(x) = \left(\frac{\partial}{\partial x_0} f(x), \dots, \frac{\partial}{\partial x_{n-1}} f(x) \right)$
 - ▶ α is a meta-parameter called “learning rate”
- ▶ Until convergence, i.e. until x remains (nearly) constant

Example: Gradient descent



Applying gradient descent to linear regression

► Objective: Minimize function $SSE(w) = \frac{1}{2} \sum_{i=0}^m (y^{(i)} - w^T x^{(i)})^2$

► Partial derivatives: $\frac{\partial}{\partial w_j} f(w) = - \sum_{i=0}^m x_j^{(i)} (y^{(i)} - w^T x^{(i)})$

► Batch update rule:

$$w_j \leftarrow w_j + \alpha \left(\sum_{i=0}^m x_j^{(i)} (y^{(i)} - w^T x^{(i)}) \right) \text{ (for every } j \text{)}$$


► Disadvantage: Requires to loop over the whole training set before anything is improved

► Alternative: “Stochastic gradient descent”

► for $i = 0, \dots, m$:

$$w_j \leftarrow w_j + \alpha \left(x_j^{(i)} (y^{(i)} - w^T x^{(i)}) \right) \text{ (for every } j \text{)}$$

Example: Linear regression with gradient descent



Remarks

- ▶ The SSE function is quadratic convex. There are no local, non-global optima, no plateaus, and no side constraints.
 - ▶ In other cases, there is no guarantee that gradient descent converges to the global minimum (can converge to local minima).
 - ▶ Gradient descent can also get stuck in plateaus.
- ▶ Learning rate is often critical:
 - ▶ Too small: Slow convergence
 - ▶ Too large: Jump over minima and potential divergence

Normal Equations (1/2)

- ▶ Gradient descent can be slow
(require many (expensive) iterations)
- ▶ Can the optimal \hat{w} be computed directly?
- ▶ For linear regression: Yes.

$$\begin{aligned}SSE(w) &= \frac{1}{2} \sum_{i=0}^m \left(y^{(i)} - w^T x^{(i)} \right)^2 \\&= \frac{1}{2} \left\| \begin{pmatrix} y^{(0)} \\ \vdots \\ y^{(m)} \end{pmatrix} - \begin{pmatrix} 1 & x_0^{(0)} & \dots & x_{n-1}^{(0)} \\ \vdots & \vdots & & \vdots \\ 1 & x_0^{(m)} & \dots & x_{n-1}^{(m)} \end{pmatrix} \begin{pmatrix} w_0 \\ \vdots \\ w_n \end{pmatrix} \right\|_2^2 \\&= \frac{1}{2} \|y - Xw\|_2^2 = \frac{1}{2} (y - Xw)^T (y - Xw)\end{aligned}$$

Normal Equations (2/2)

$SSE(w)$ is quadratic and ≥ 0 for all w . Thus, its global minimum is the unique root of its derivative:

$$\begin{aligned} 0 &= \nabla_w SSE(\hat{w}) \\ &= \frac{1}{2} \nabla_w \left[(y - Xw)^T (y - Xw) \right] \Big|_{w=\hat{w}} \\ &= \frac{1}{2} \nabla_w \left[y^T y - w^T X^T y - y^T Xw + w^T X^T Xw \right] \Big|_{w=\hat{w}} \\ &= -X^T y + X^T X \hat{w} \end{aligned}$$

Since X (usually) has full column rank:

$$\hat{w} = (X^T X)^{-1} X^T y$$

Online Learning

- ▶ $\hat{w} = (X^T X)^{-1} X^T y$
- ▶ Can this be calculated on a mobile system?
- ▶ Memory/processing effort with new incoming (Big) data?
- ▶ Sherman-Morrison-Woodbury formula
 - Fast online update for $(X(m+1)^T X(m+1))^{-1}$
 - ▶ $X(m+1)^T X(m+1) = X(m)^T X(m) + (x^{(m+1)})^T x^{(m+1)}$
 - ▶ Fixed dimensions of $(X(m)^T X(m))^{-1}$ and $X(m)^T y(m)$
 - ▶ $X(m+1)^T y(m+1) = X(m)^T y(m) + (x^{(m+1)})^T y^{(m+1)}$

Other target functions

- ▶ The “linear function plus gaussian noise” model is often too simple.
- ▶ More general classes of functions have a better chance of fitting the data well
- ▶ Classes of functions can be defined by a projection function ϕ
- ▶ Examples of classes of functions (for $X = \mathbb{R}$):
 - ▶ Linear functions: $\phi(x) = (1, x)^T$
 - ▶ Polynomials: $\phi(x) = (x^0, x^1, \dots, x^k)^T$, k = degree of polynomial
 - ▶ Linear combination of a set of sinusoidal functions:
 $\phi(x) = (\sin(0x/k), \sin(1x/k), \sin(2x/k), \dots, \sin(x))^T$

Other target functions

Objective: Minimize function


$$\begin{aligned}SSE(w) &= \frac{1}{2} \sum_{i=0}^m \left(y^{(i)} - w^T \phi(x^{(i)}) \right)^2 \\ &= \frac{1}{2} (y - X_\phi w)^T (y - X_\phi w)\end{aligned}$$

$$\text{with } X_\phi = \begin{pmatrix} \phi(x^{(0)})_0 & \phi(x^{(0)})_1 & \dots & \phi(x^{(0)})_k \\ \vdots & \vdots & & \vdots \\ \phi(x^{(m)})_0 & \phi(x^{(m)})_1 & \dots & \phi(x^{(m)})_k \end{pmatrix}$$

This can still be solved with the normal equations:

$$\hat{w} = (X_\phi^T X_\phi)^{-1} X_\phi^T y$$

Gradient descent for $\phi(x) = (x^0, x^1, x^2, x^3)$




How to select ϕ

- ▶ The function ϕ defines which functions can be represented. It is also called “the model”
- ▶ Intuitively, ϕ should somehow reflect the properties of the unknown target function.
- ▶ But even when it is known that the true function is a polynomial: How should k be chosen?
- ▶ First idea:
 - ▶ Choose k as large as possible since the class of learnable functions becomes bigger.
 - ▶ Even linear functions remain learnable
(set all higher-order w_i to 0)
- ▶ Is this really a good idea?


Polynomial model

Degree=0




Polynomial model

Degree=1




Polynomial model

Degree=2




Polynomial model

Degree=3




Polynomial model

Degree=4




Polynomial model

Degree=5




Polynomial model

Degree=6




Polynomial model

Degree=7





Polynomial model

Degree=8



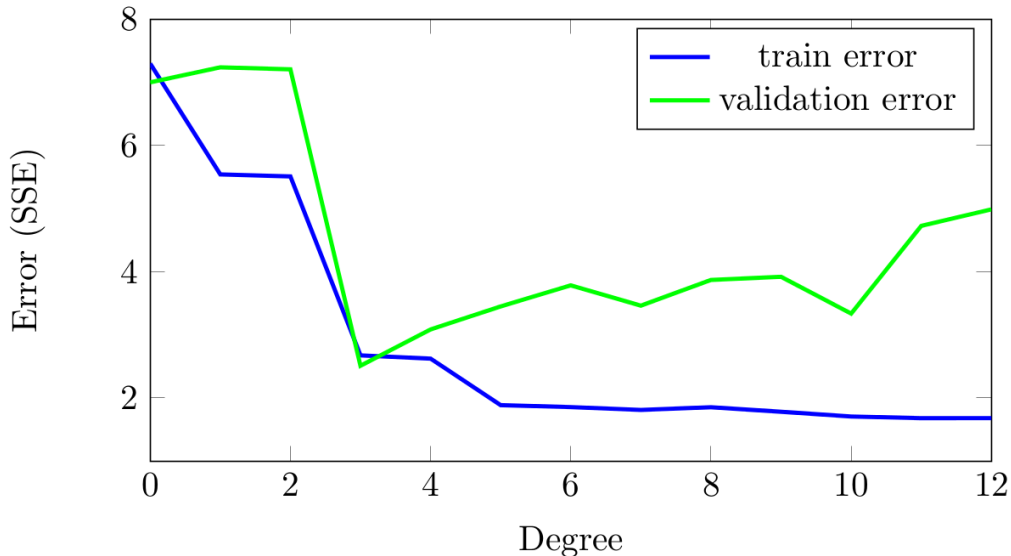
Train- and Testerror



Overfitting and Underfitting

- ▶ Larger k allows to fit the training data points better
- ▶ But: The hypothesis was maximally similar to true function for $k = 3$.
- ▶ Explanation: Large values of k (an “increased complexity”) allows to fit to the noise in the training data (“Overfitting”)
- ▶ Choosing k too small results in a overly simple hypothesis (“Underfitting”)
- ▶ How can k be chosen when true function is unknown?
- ▶ One approach: Test learned hypothesis on hold-out validation data (\rightarrow cross-validation)

Train- and Validation-Error



Model Selection

- ▶ This approach can be (partially) automated!
- ▶ Split labeled data \mathcal{D} into \mathcal{D}_{train} and $\mathcal{D}_{validation}$
- ▶ For each value of k :
 - ▶ Determine a hypothesis on \mathcal{D}_{train} (using normal equations)
 - ▶ Evaluate learned hypothesis on $\mathcal{D}_{validation}$
- ▶ Select k with lowest SSE on validation data
- ▶ Train a polynomial with degree k on whole \mathcal{D}
- ▶ Can be made more accurate (cross-validation) and more general (not restricted to polynomial models)

Advanced Topics

Regularization

- ▶ Techniques like cross-validation allow to choose appropriate model complexity
- ▶ But: CV is computationally costly (besides some other minor drawbacks)
- ▶ Alternative: Regularization
- ▶ Idea behind regularization: Use model with high complexity but enforce that solution is somehow “smooth”
- ▶ Regularization is also recommendable from a computational point of view: It reduces problems that occur due to numerical instability.


Tikhonov Regularization

- ▶ Linear Least Squares: $\hat{w} = \arg \min_w \frac{1}{2} \|y - Xw\|_2^2$
- ▶ Linear Least Squares + Tikhonov Regularization:
$$\hat{w} = \arg \min_w \frac{1}{2} \|y - Xw\|_2^2 + \|\Gamma w\|_2^2$$
- ▶ In many cases, Γ is chosen to be diagonal, e.g. $\Gamma = \text{diag}(\tau)$ for some $\tau \in \mathbb{R}$. Could set bias to certain weights.
- ▶ For $\tau = 0$: Equal to (unregularized) Linear Least Squares
- ▶ For $\tau > 0$: Giving preference to solutions w with smaller norms (the stronger the larger τ)
- ▶ Normal equation solution:

$$\hat{w} = (X^T X + \Gamma^T \Gamma)^{-1} X^T y$$


Ridge regression example

$\text{Tau}=0.0$, Degree=9




Ridge regression example

$\text{Tau}=0.2$, Degree=9




Ridge regression example

$\text{Tau}=0.5$, Degree=9




Ridge regression example

$\text{Tau}=1.0$, Degree=9




Ridge regression example

$\text{Tau}=1.5$, Degree=9




Ridge regression example

$\text{Tau}=2.0$, Degree=9




Ridge regression example

$\text{Tau}=3.0$, Degree=9




Ridge regression example

$\text{Tau}=5.0$, Degree=9




Ridge regression example

$\text{Tau}=10.0$, Degree=9




Ridge regression example

$\text{Tau}=25.0$, Degree=9




Ridge regression example

$\text{Tau}=50.0$, Degree=9





Ridge regression example


$\text{Tau}=75.0, \text{Degree}=9$



Example: SSE on Train and Test Data



Example: Coefficients (k=9)



Left: $\tau = 0$; right: true coefficients (blue==dark blue)

Remarks

- ▶ Linear regression combined with Tikhonov regularization is also called “Ridge Regression”
- ▶ Regularization cannot magically remove the risk of overfitting
- ▶ Instead of choosing the model complexity one has to choose the regularization parameter(s)
- ▶ But: Often the specific choice of the regularization parameter is less critical
- ▶ Furthermore: Remedies issues with numeric stability
- ▶ Even though Tikhonov regularization seems rather arbitrary, it can be motivated statistically
(prior, German: A-priori-Wahrscheinlichkeit)

Thank You!

Please feel free to ask questions in the forums.