

# META LEARNING

Machine Learning for Autonomous Robots

Dr. Matias Valdenegro-Toro

Slides by Anett Seeland, Mario Michael Krell, Alexander Fabisch

Assistant Professor at University of Groningen

December 13, 2022 – Bremen, Deutschland

- 1 Overview of Meta Learning
- 2 Introduction to Ensemble methods
- 3 Bias-Variance Trade-off
- 4 Bagging
- 5 Boosting
- 6 Modern Ensemble Methods
- 7 references

## Overview of Meta Learning

# What is Meta Learning?

Meta-learning in this lecture is different from the *Meta-Learning* in the current Deep Learning literature.

- ▶ Meta-Learning here refers to classic models of models (the meta part).
- ▶ Meta-Learning in Deep Learning is also models of models but more related to things like *learning to learn* and its use with neural network models.
- ▶ This is just a word of warning in case you perform a literature search and are a bit confused...

# What is Meta Learning?

Meta-learning includes... (ordered by degree of automization)

- ▶ guidelines and heuristics for the application of algorithms
- ▶ **tools that enable users without experience in machine learning (ML) to solve tasks with ML, these tools include**
  - ▶ broad set of algorithms
  - ▶ evaluation methods
  - ▶ model selection
- ▶ meta-level learning schemes:
  - ▶ multi-task learning
  - ▶ transfer learning
  - ▶ lifelong learning
  - ▶ **ensemble learning**
- ▶ automatic machine learning (*active field of research*)


Why not just one algorithm that solves everything?

## There is no free lunch.

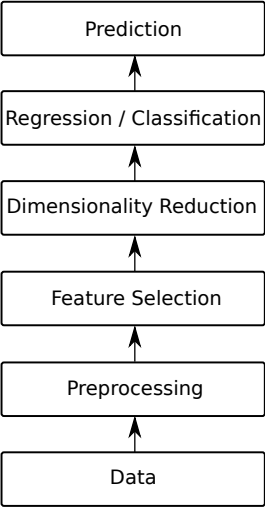
- ▶ Every learning algorithm has an **inductive bias**.
- ▶ No learning algorithm is better than any other learning algorithm **on average over all datasets**.
- ▶ No Free Lunch (NFL) theorem.

## Inductive Bias

- ▶ An unbiased algorithm cannot generalize. Having no inductive bias means immediate overfitting.
- ▶ The bias of a learning algorithm is called **inductive bias**.
- ▶ Examples: min. CV error, maximum margin, Occam's razor

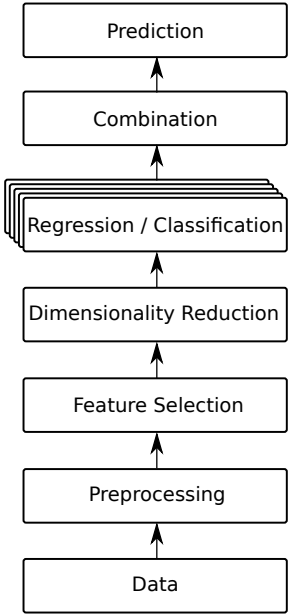


# Standard Pipeline






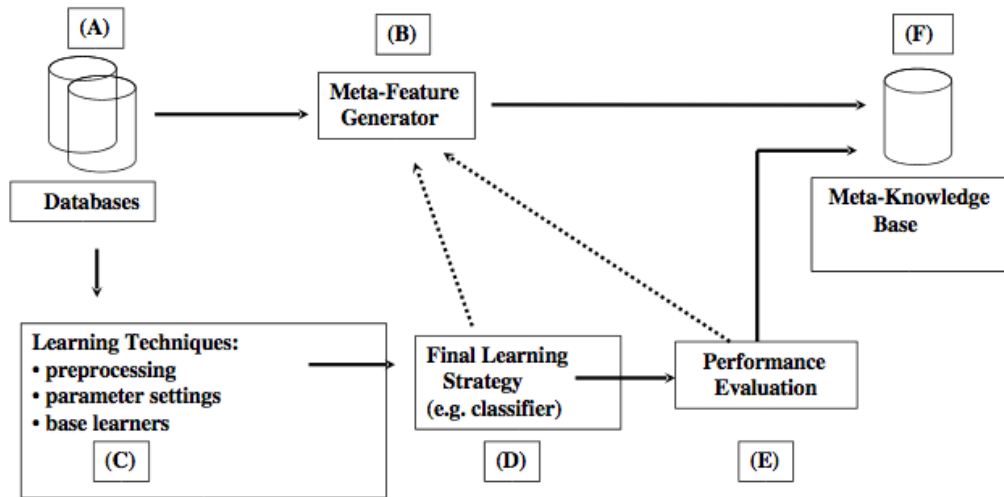
# Ensemble Learning



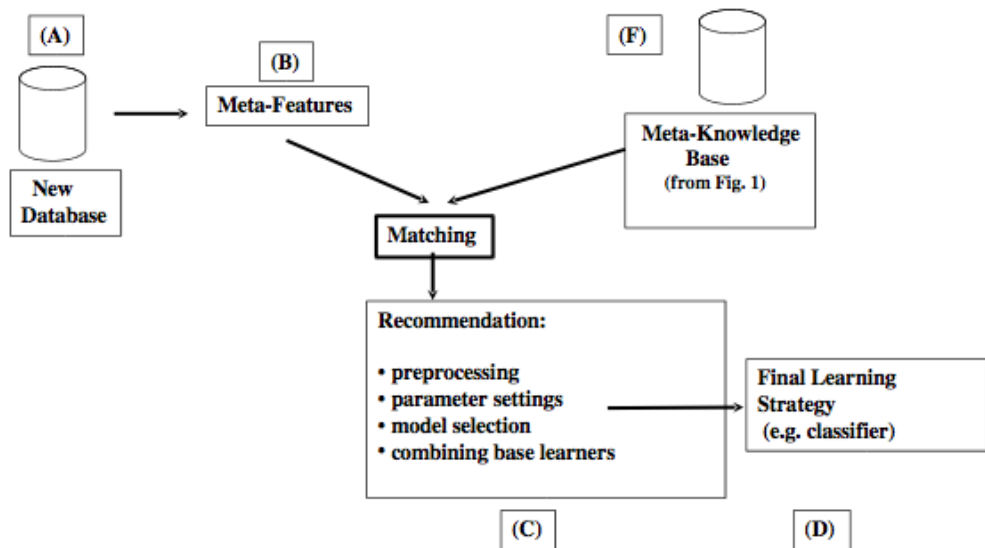
# Outlook



## The Knowledge-Acquisition Mode



## The Advisory Mode



# Techniques in Meta Learning

## Dataset Characterization:

- ▶ Statistical and Information-Theoretic Characterization
  - ▶ # classes, # features, # examples,  $\frac{\# \text{ examples}}{\# \text{ features}}$
  - ▶ degree of correlation between features and target concept
  - ▶ skewness, kurtosis, signal-to-noise ratio
- ▶ Model-Based Characterization
  - ▶ e.g. support vector machine: C, kernel, kernel parameters
- ▶ Landmarking
  - ▶ performance of different learning mechanisms
  - ▶ sampling landmarks

## Introduction to Ensemble methods

## Motivation

When *wise people* make **critical** decisions,  
they usually take into account  
the opinions of *several* experts  
rather than relying  
on their own judgment or  
that of a solitary trusted adviser.

## Ensemble methods

Ensemble: using many classifiers together

- ▶ could be different types of classifiers
- ▶ each classifier gives a different view
- ▶ usually each is trained differently
- ▶ need a sensible method of combining classifiers

Popular methods:

- ▶ bagging
- ▶ boosting
- ▶ stacking
- ▶ random forest (see Classification 1)



## Combining classifiers

Suppose that we have:

- ▶  $m$  labeled examples  $\langle \mathbf{x}, f(\mathbf{x}) \rangle$ , where  $f : R \rightarrow \{-1, +1\}$
- ▶ a number of  $T$  hypotheses  $h_1, \dots, h_T$ , each trained on the data
- ▶ each  $h_t : R \rightarrow \{-1, +1\}$

How should we combine the  $T$  classifiers?

- ▶ choose  $h_t$  with the least training error
- ▶ choose  $h_t$  with the least validation error
- ▶ use all  $T$  classifiers in a sum
- ▶ use all  $T$  classifiers in a weighted sum

## Bias-Variance Trade-off

# Bias-Variance Decomposition

In the **ideal situation** of an infinite dataset:

- error will still occur because no learning scheme is perfect
- **bias** for the learning problem

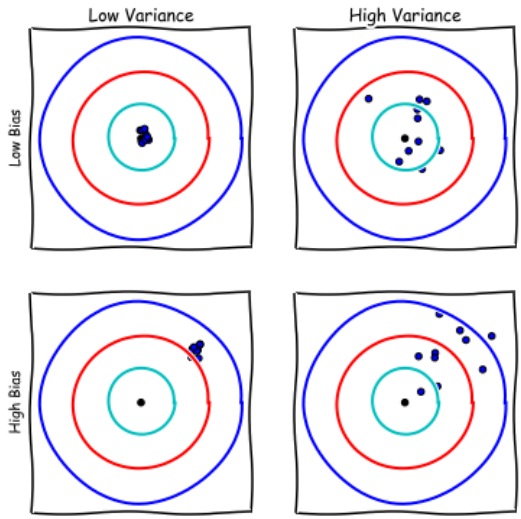
In **practical situations** there is another error:

- ▶ stems from the particular training set used (unavoidably finite)
- not fully representative of the actual population of instances
- **variance** for the learning method

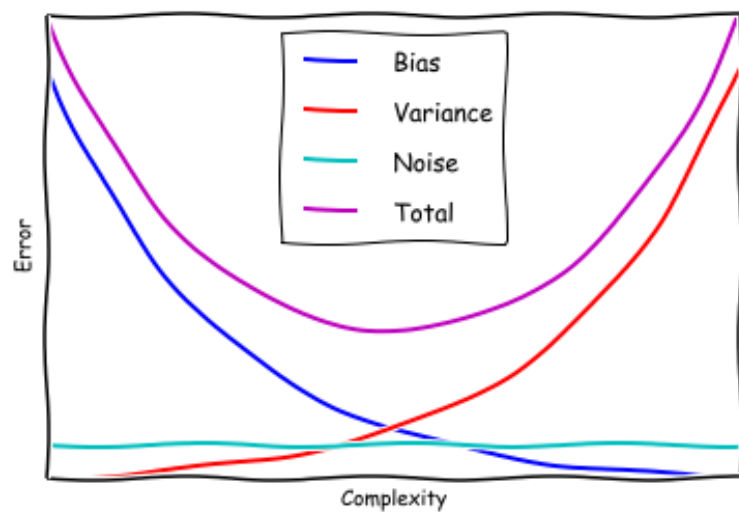
total expected error is a sum of **bias** and **variance** (and noise)

- ⇒ combining multiple classifiers decreases the expected error by reducing the variance component

# Bias / Variance



Bias / Variance



## Inductive Bias vs. Bias Error

Can you imagine an example where we have an inductive bias but no bias error?

### **Example:**

- ▶ We know that we want to approximate a sine function:

$$f(x) = a \sin(bx + c)$$

- ▶ We don't know  $a, b, c$
- ▶ We can fit  $a, b, c$  from data
- Strong inductive bias
- No bias error

Bagging

# Unstable learners

Imagine:

- ▶ several randomly chosen training sets of the same size
- ▶ build a decision tree for each dataset

Will the trees make the same predictions?

- No! (particularly for small datasets)
- reason: unstable learning process
- ⇒ voting becomes more reliable with more votes
- ⇒ combined classifiers will seldom be less accurate
- ⇒ but improvement is not guaranteed



## Bagging – bootstrap aggregation

- ▶ reduces variance (typically uses complex models)
- ▶ random sampling with replacement from dataset
- ▶ applies base learner to each new dataset
- ▶ votes with equal weights

combined model ...

- + ... often performs significantly better than single model
- + ... is never substantially worse

## Bagging – bootstrap aggregation

```
Bagging( $\mathcal{D}, T$ ):  
  model generation( $\mathcal{D}, T$ ):  
     $m \leftarrow |\mathcal{D}|$   
    for  $t \in \{1, \dots, T\}$  do  
       $\mathcal{D}_t \leftarrow \text{bootstrap}(\mathcal{D}, m)$   
       $h_t \leftarrow \text{build\_model}(\mathcal{D}_t)$   
    end  
  classification( $x$ ):  
    for  $t \in \{1, \dots, T\}$  do  
       $\text{store}(h_t(x))$   
    end  
  return class that has been predicted most often
```

Boosting

## Ideas behind Boosting

- ▶ reduces bias (typically uses simple models)
- ▶ design significantly different models
- ▶ each model treats a reasonable percentage of the data correctly
- aim: models complement one another, each being a specialist in a part of the domain
- ▶ weighting is used to give more influence to the more successful models

# Bagging vs. Boosting

## Similarities:

- ▶ both use voting
- ▶ both combine models of the same type, e.g. decision trees

## Differences:

- ▶ Boosting is sequential
- ▶ Boosting weights a model's contribution by its performance
- ▶ Boosting reduces bias, bagging reduces variance

# The AdaBoost Algorithm

AdaBoost( $\mathcal{D}, T$ ):

**model generation**( $\mathcal{D}, T$ ):

**for**  $i \in \{1, \dots, m\}$  **do**  $W_1(i) \leftarrow 1/m$  // sample weights

        ;

**for**  $t \in \{1, \dots, T\}$  **do**

$h_t \leftarrow \text{build\_model}(\mathcal{D}, W_t)$ ;

$\epsilon_t \leftarrow \text{error}(\mathcal{D}, W_t, h_t)$ ;

**if**  $\epsilon_t = 0$  or  $\epsilon_t \geq 1/2$  **then** continue;

$\alpha_t \leftarrow \frac{1}{2} \ln\left(\frac{1-\epsilon_t}{\epsilon_t}\right)$ ;

**for**  $i \in \{1, \dots, m\}$  **do**

$W_{t+1}(i) \leftarrow \frac{W_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = f(x_i) \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq f(x_i) \end{cases}$  ;

                //  $Z_t$  to normalize  $W_{t+1}$

**end**

**end**

# The AdaBoost Algorithm

AdaBoost( $\mathcal{D}, T$ ):

$\vdots$

**classification( $x$ ):**

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

If model class can handle weighted instances:

$$\epsilon = \frac{\sum w_i}{\sum w_j} \text{ for } i \in \{i | h(x_i) \neq f(x_i)\} \text{ and } j \in \{1, \dots, m\}$$

Otherwise, resample the data according to a distribution  $W_t$

How much can the weights change?

$$W_{t+1}(i) \leftarrow \frac{W_t(i)}{Z_t} \times \begin{cases} e^{-\alpha_t} & \text{if } h_t(x_i) = f(x_i) \\ e^{\alpha_t} & \text{if } h_t(x_i) \neq f(x_i) \end{cases}$$

recall:

- ▶  $0 < \epsilon_t < 0.5$
- ▶  $\alpha_t = \frac{1}{2} \ln\left(\frac{1-\epsilon_t}{\epsilon_t}\right) \rightarrow \alpha \in (0, \infty)$
- ▶  $e^{-\alpha_t} = \sqrt{\epsilon_t/(1-\epsilon_t)}$ , correct  $\rightarrow$  sample weight decreases
- ▶  $e^{\alpha_t} = \sqrt{(1-\epsilon_t)/\epsilon_t}$ , error  $\rightarrow$  sample weight increases

Thus, if  $\epsilon_t \approx 0$  then  $W_t(i) \rightarrow 0$



## Role of $\alpha_t$

Recall the final classifiers's form:

$$H(x) = \text{sign} \left( \sum_{t=1}^T \alpha_t h_t(x) \right)$$

Also,

$$\alpha_t = \frac{1}{2} \ln \left( \frac{1 - \epsilon_t}{\epsilon_t} \right)$$

Thus,

- ▶  $\epsilon_t = 0.5 \rightarrow \alpha_t = 0$
- ▶  $\epsilon_t < 0.5 \rightarrow \alpha_t > 0$
- ▶  $\epsilon_t > 0.5 \rightarrow \alpha_t < 0$ , will not happen, because those classifiers are ignored!

## Greedy search

Interestingly, AdaBoost is a greedy strategy.

It always minimizes the error with respect to the current  $W_t$ , and does not backtrack.

## How does boosting help?


### Assumption:

- ▶ easy to find *weak learners* that are “often” correct (slightly more than random)
- ▶ hard to find a *strong learner*, that is highly accurate

### Combination of weak learners to generate a strong learner:


- ▶ focuses effort on hard-to-classify examples
- ▶ takes hard work off the learner
- ▶ learners only have to specialize in small areas
- ▶ can identify outliers

How does boosting help?



# Gradient boosting

- ▶ generalizes AdaBoost to optimize an arbitrary differentiable loss function
- inspired from regression



# Gradient boosting

- ▶ boosting can be seen as an iterative functional gradient descent algorithm
- ▶ example: MSE for regression
- residuals are the negative gradients of the squared error loss function

$$-\nabla_{F(x)} \left[ \frac{1}{2} (y - F(x))^2 \right] = y - F(x)$$


## Gradient Boosting – Idea of the algorithm

- ▶ input:  $\mathcal{D}$ ,  $T$ , and some specified loss function  $L(y, F(x))$
- ▶ in each step, the previous model  $F_{t-1}$  is improved by adding an estimator  $h$  (weak learner)
- $F_t = F_{t-1} + h(x)$
- ▶ How to find  $h$ ?
- optimal  $h$  would imply

$$F_t = F_{t-1} + h(x) = y \Leftrightarrow h(x) = y - F_{t-1}$$

- train  $h$  that approximates the negative gradient of  $L$
- ▶  $F_t(x) = F_{t-1}(x) + \gamma_t h_t(x)$
- ▶  $\gamma_t = \arg \min_{\gamma} \sum_{i=1}^m L(y_i, F_{t-1}(x_i) + \gamma h_t(x_i))$

# Gradient boosting for decision tree regression





## Modern Ensemble Methods

# Deep Ensembles

- ▶ Simple idea, make an ensemble of neural networks [Lakshminarayanan et al. 2017].
- ▶ Trained on the same data, models converge to different solutions due to random weight initialization.
- ▶ It has properties of good uncertainty quantification and out of distribution detection.
- ▶ It can estimate aleatoric (data) and epistemic (model) uncertainty. For aleatoric uncertainty in regression it uses the following loss:

$$-\log p_{\theta}(y_n|x_n) = 0.5 \left( \log \sigma_{\theta}^2(x) + \frac{(y - \mu_{\theta}(x))^2}{\sigma_{\theta}^2(x)} \right)$$

# Deep Ensembles

## Classification

Take average of ensemble member output probabilities.


## Regression

Each ensemble member outputs  $\mu_{\theta_m}(x)$  and  $\sigma_{\theta_m}^2(x)$ , and they are combined as a mixture of gaussians  $M^{-1} \sum \mathcal{N}(\mu_{\theta_m}(x), \sigma_{\theta_m}^2(x))$  represented as  $\mathcal{N}(\mu_*(x), \sigma_*^2(x))$  where:


$$\mu_*(x) = M^{-1} \sum_m \mu_{\theta_m}(x)$$

$$\sigma_*^2(x) = M^{-1} \sum_m (\sigma_{\theta_m}^2(x) + \mu_{\theta_m}^2(x)) - \mu_*^2(x)$$

# Deep Ensembles - Performance



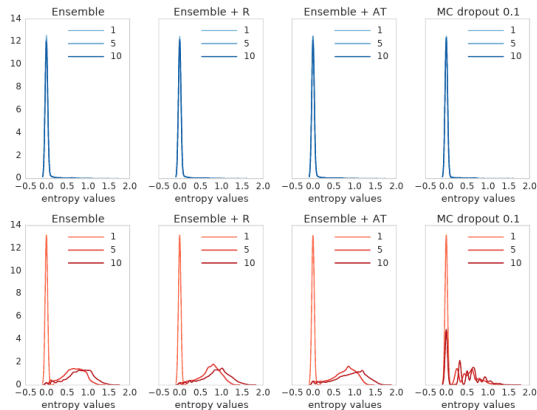
(a) MNIST dataset using 3-layer MLP



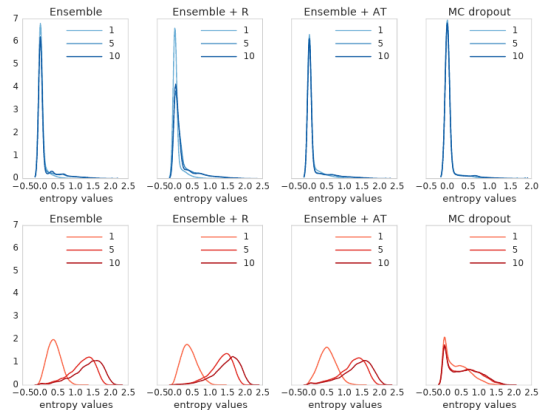
(b) SVHN using VGG-style convnet

Figure 2: Evaluating predictive uncertainty as a function of ensemble size  $M$  (number of networks in the ensemble or the number of MC-dropout samples): Ensemble variants significantly outperform MC-dropout performance with the corresponding  $M$  in terms of all 3 metrics. Adversarial training improves results for MNIST for all  $M$  and SVHN when  $M = 1$ , but the effect drops as  $M$  increases.

# Deep Ensembles - Out of Distribution Detection



(a) MNIST-NotMNIST



(b) SVHN-CIFAR10

references

## Literature

- ▶ “Meta-Learning – Concepts and Techniques”, O. Maimon, L. Rokach (eds.), In: *Data Mining and Knowledge Discovery Handbook* (2nd Ed.), 2010
- ▶ “Data Mining: Practical Machine Learning Tools and Techniques” (2nd Ed.), Ian H. Witten, Eibe Frank, Morgan Kaufmann, 2005
- ▶ “A short introduction to boosting”, Y. Freund, R. Schapire, N. Abe, In: JOURNAL-JAPANESE SOCIETY FOR ARTIFICIAL INTELLIGENCE, Vol. 14, p.771–780, 1999
- ▶ Lakshminarayanan B, Pritzel A, Blundell C. Simple and scalable predictive uncertainty estimation using deep ensembles. NeurIPS 2017.

Thank You!  
Please feel free to ask questions in the  
forums.